PS5: DNA Sequence Alignment

Due: Monday, April 1, 11:59pm

The goals of this assignment are to solve a fundamental problem in computational biology and to learn about a powerful programming paradigm known as *dynamic programming*.

On this assignment, you are encouraged (but not required) to work with a partner provided you practice **pair programming**. Pair programming is "*a practice in which two programmers work side-by-side at one computer, continuously collaborating on the same design, algorithm, code, or test*." One partner is driving (designing and typing the code) while the other is navigating (reviewing the work, identifying bugs, and asking questions). The two partners switch roles every 30-40 minutes, and on demand, brainstorm.

If you are working with a partner, you should remember

- Make a single submission with both partners names listed in the group when submitting to Gradescope.

- You are responsible for making sure that you were included in the submission.

- You are responsible for making sure that your submission fulfills the project requirements.

- You should clearly identify your partner in your `Readme-ps5.md` file.

For this project, we will:

- Write a program to compute the optimal sequence alignment of two DNA strings. This program will introduce you to the field of computational biology in which computers are used to do research on biological systems. Further, you will be introduced to a powerful algorithmic design paradigm known as dynamic programming.

- Use `valgrind`, a memory analysis tool.

- Measure and report the space and time performance of the implementation.

Average time to complete assignment: $\sim$ 5 hours.

1 Understanding the Problem

1.1 Biology Review

A genetic sequence is a string formed from a four-letter alphabet (Adenine (A), Thymine (T), Guanine(G), and Cytosine (C)) of biological macromolecules referred to together as the DNA bases. A gene is a genetic sequence that contains the information needed to construct a protein. All of your genes taken together are referred to as the human genome, a blueprint for the parts needed to construct the proteins that form your cells. Each new cell produced by your body receives a copy of the genome. This copying process, as well as the natural wear and tear, introduces a small number of changes into the sequences of many genes. Among the most common changes are the substitution of one base for another and the deletion of a substring of bases; such changes are generally referred to as point mutations. As a result of these point mutations, the same gene sequence from closely related organisms will have slight differences.

1.2 The Problem

Through your research, you have found the following sequence of a gene in a previously unstudied organism.

`A A C A G T T A C C`

What is the function of the protein that this gene encodes? You could begin a series of uninformed experiments in the lab to determine what role this gene plays. However, there is a good chance that it is a variant of a known gene in a previously studied organism. Since biologists and computer scientists have laboriously determined (and published) the genetic sequences of many organisms (including humans), you would like to leverage this information to your advantage. We'll compare the above genetic sequence with one which has already been sequenced and whose function is well understood.

`T A A G G T C A`

If the two genetic sequences are similar enough, we might expect them to have similar functions. We would like a way to quantify "similar enough".

1.3 Edit Distance

In this assignment, we will measure the similarity of two genetic sequences by their edit distance, a concept first introduced in the context of coding theory, but which is now widely used in spell checking, speech recognition, plagiarism detection, file revisioning, and computational linguistics. We align the two sequences, but we are permitted to insert gaps in either sequence (e.g. to make them have the same length). We pay a penalty for each gap that we insert and also for each pair of characters that mismatch in the final alignment. Intuitively, these penalties model the relative likeliness of point mutations arising from insertion/deletion and substitution. We produce a numerical score according to the following table, which is widely used in biological applications.

Insert a gap	2
Align two characters that mismatch	1
Align two characters that match	0

Here are two possible alignments of the strings `x = "AACAGTTACC"` and `y = "TAAGGTCA"`.

x	y	cost
A	T	1
A	A	0
C	A	1
A	G	1
G	G	0
T	T	0
T	C	1
A	A	0
C	-	2
C	-	2
		8

x	y	cost
A	T	1
A	A	0
C	-	2
A	A	0
G	G	0
T	G	1
T	T	0
A	-	2
C	C	0
C	A	1
		7

The first alignment has a score of 8 while the second one has a score of 7. The edit-distance is the score of the best possible alignment between the two genetic sequences over all possible alignments. In this example, the second alignment is in fact optimal, so the edit-distance between the two strings is 7. Computing the edit-distance is a nontrivial computational problem because we must find the best alignment among exponentially many possibilities. For example, if both strings are 100 characters long, then there are more than $10^7 5$ possible alignments.

A recursive solution is an elegant approach. However, it is far too inefficient because it recalculates each sub-problem over and over. Once we have defined the recursive definition we can redefine the solution using a dynamic programming approach which calculates each sub-problem once.

1.4 A recursive solution

We will calculate the edit-distance between the two original strings x and y by solving the edit-distance problem on smaller suffixes of the two strings. We use the notation $x[i]$ to refer to character i of the string. We also use the notation $x[i : M]$ to refer to the substring of x consisting of the characters $x[i], x[i + 1], \ldots, x[M - 1]$. Finally, we use the notation $opt[i][j]$ to denote the edit distance of $x[i : M]$ and $y[j : N]$. For example, consider the two strings x = "AACAGTTACC" and y = "TAAGGTCA" of length $M = 10$ and $N = 8$, respectively. Then $x[2]$ is 'C' and $x[2 : M]$ is "CAGTTACC" and $y[8 : N]$ is the empty string. The edit distance of x and y is $opt[0][0]$.

Now we describe the recursive scheme for computing the edit distance of $x[i : M]$ and $y[j : N]$. Consider the pair of characters in an optimal alignment of $x[i : M]$ with $y[j : N]$. There are three possibilities:

1. The optimal alignment matches $x[i]$ up with $y[j]$. In this case, we pay a penalty of either 0 or 1, depending on whether $x[i]$ equals $y[j]$, plus we still need to align $x[i + 1 : M]$ with $y[j + 1 : N]$. What is the best way to do this? This sub-problem is exactly the same as the original sequence alignment problem, except that the two inputs are each suffixes of the original inputs. Using our notation, this quantity is $opt[i + 1][j + 1]$.

2. The optimal alignment matches $x[i]$ up with a gap. In this case, we pay a penalty of 2 for a gap and still need to align $x[i + 1 : M]$ with $y[j : N]$. This sub-problem is identical to the original sequence alignment problem except that the first input is a proper suffix of the original input.

3. The optimal alignment matches $y[j]$ up with a gap. In this case, we pay a penalty of 2 for a gap and still need to align $x[i : M]$ with $y[j + 1 : N]$. This sub-problem is identical to the original sequence alignment problem except that the first input is a proper suffix of the original input.

The key observation is that all of the resulting sub-problems are sequence alignment problems on suffixes of the original inputs. To summarize, we can compute $opt[i][j]$ by taking the minimum of three quantities:

```
opt[i][j] = min{ opt[i+1][j+1] + 0/1, opt[i+1][j] + 2, opt[i][j+1] + 2 }
```

This equation works assuming $i < M$ and $j < N$. Aligning an empty string with another string of length k requires inserting k gaps, for a total cost of $2k$. Thus, in general, we should set $opt[M][j] = 2(N - j)$ and $opt[i][N] = 2(M - i)$. For our example, the final matrix is:

		0	1	2	3	4	5	6	7	8
		T	A	A	G	G	T	C	A	-
0	A	7	8	10	12	13	15	16	18	20
1	A	6	6	8	10	11	13	14	16	18
2	C	6	5	6	8	9	11	12	14	16
3	A	7	5	4	6	7	9	11	12	14
4	G	9	7	5	4	5	7	9	10	12
5	T	8	8	6	4	4	5	7	8	10
6	T	9	8	7	5	3	3	5	6	8
7	A	11	9	7	6	4	2	3	4	6
8	C	13	11	9	7	5	3	1	3	4
9	C	14	12	10	8	6	4	2	1	2
10	-	16	14	12	10	8	6	4	2	0

By examining $opt[0][0]$, we conclude that the edit distance of x and y is 7.

1.5 A dynamic programming approach

A direct implementation of the above recursive scheme will work, but it is spectacularly inefficient. If both input strings have n characters, than the number of recursive calls will exceed 2^n. To overcome this performance bug, we use *dynamic programming*. Dynamic programming is a powerful algorithmic paradigm, first introduced by Bellman in the context of operations research and then applied to the alignment of biological sequences by Needlman

and Wunsch. Dynamic programming now plays the leading role in many computational problems, including control theory, financial engineering, and bioinformatics, including BLAST (the sequence alignment program almost universally used by molecular biologists in their experimental work). The key idea of dynamic programming is to break up a large computational problem into smaller sub-problems, store the answers to those smaller sub-problems, and, eventually, use the stored answers to solve the original problem. This avoids recomputing the same quantity over and over again. Instead of using recursion, use a nested loop that calculates $opt[i][j]$ in the right order so that $opt[i+1][j+1]$, $opt[i+1][j]$, and $opt[i][j+1]$ are all computed before we try to compute $opt[i][j]$.

1.6   Recovering the alignment itself

The above procedure describes how to compute the edit distance between two strings. We now outline how to recover the optimal string alignment itself. The idea is to retrace the steps of the dynamic programming algorithm backwards, re-discovering the path of choices (highlighted in red in the table above) from $opt[0][0]$ to $opt[M][N]$. To determine the choice that led to $opt[i][j]$, we consider the three possibilities:

1. The optimal alignment matches $x[i]$ up with $y[j]$. In this case, we must have $opt[i][j] = opt[i+1][j+1]$ if $x[i]$ equals $y[j]$ or $opt[i][j] = opt[i+1][j+1] + 1$ otherwise.

2. The optimal alignment matches $x[i]$ up with a gap. In this case, we must have $opt[i][j] = opt[i+1][j] + 2$.

3. The optimal alignment matches $y[j]$ up with a gap. In this case, we must have $opt[i][j] = opt[i][j+1] + 2$.

Depending on which of the three cases apply, we move diagonally, down, or right towards $opt[M][N]$, printing out $x[i]$ aligned with $y[j]$ (case 1), $x[i]$ aligned with a gap (case 2), or $y[j]$ aligned with a gap (case 3). In the example above, we know that the first T aligns with the first A because $opt[0][0] = opt[1][1]+1$, but $opt[0][0] \neq opt[1][0]+2$ and $opt[0][0] \neq opt[0][1]+2$.

The optimal alignment is:

x	y	cost
A	T	1
A	A	0
C	-	2
A	A	0
G	G	0
T	G	1
T	T	0
A	-	2
C	C	0
C	A	1

2   Implementation

You may elect to implement any of four solution approaches:

- Recursive without memoization (two slow to be practical, but order n in space).

- Recursive using memoization.

- Dynamic programming using an $n \times m$ matrix using the Needleman-Wunsch method

- Using Hirschberg's algorithm, which is linear in space.

If you implement the dynamic programming with a matrix approach, remember that the solution is in two parts:

1. Filling out the $n \times m$ matrix per the min-of-three-options formula, bottom to top, right to left (this gives you the optimal edit distance in the upper-left ([0][0]) cell of the matrix)

2. Traversing the matrix from top-to-bottom, left-to-right (i.e. [0][0] to [n][m]) to recover the choices you made in filling it, and thereby also recovering the actual edit sequence.

2.1 API Specification

You should create a class called `EDistance` (for "Edit Distance") with the following methods:

- A **constructor** that accepts the two strings to be compared, and that allocates any data structures necessary in order to do the work (e.g. the $n \times m$ matrix)

- A **static** method `int penalty(char a, char b)` that returns the penalty for aligning chars a and b (this will be either a 0 or a 1).

- A **static** method `int min3(int a, int b, int c)` which returns the minimum of the three arguments.

- A method `int optDistance()` which populates the matrix based on having the two strings, and returns the optimal distance (from the [0][0] cell of the matrix when done).

- A methoc `string alignment()` which traces the matrix and returns a string that can be printed to display the actual alignment. In general, this will be a multi-line string - i.e., with embedded \ns.

You should have a `main` routine that accepts two strings from `stdin`, uses your `EDistance` class to do the work, and then prints the result to `stdout`. Remember that your final output should look like this:

```
% ./EDistance < example10.txt
Edit distance = 7
A T 1
A A 0
C - 2
A A 0
G G 0
T G 1
T T 0
A - 2
C C 0
C A 1
```

2.2 Implementation

You have to allocate the memory for the opt matrix dynamically, after you read in the two strings and figure out how long they are. There are a few different ways to do this:

- `vector` of columns, each containing a row `vector`

- one long `vector`, with internal calculations to treat it as a matrix

- one big block of memory, similarly with calculations to treat it as a matrix

- others?

See `http://stackoverflow.com/questions/936687/how-do-i-declare-a-2d-array-in-c-using-new` for some ideas.

If you use `new`, make sure to de-allocate memory in your class destructor; e.g. if your array pointer is named `_opt`, then `delete [] _opt;`

The dynamic programming solution we discussed requires filling the whole matrix with values (step 1 of the two-part solution), so it is $O(n^2)$ in space (or more precisely $O(n \cdot m)$). You shouldn't expect to have your code work for the test case with two $500,000$ `char` strings. Assuming you use a 32-bit `int` to hold edit distance values in your matrix, that's 2 MB of data squared or 4,000 GB. Your computer probably can't allocate this much RAM.

The largest problem you should be able to handle is `ecoli28284.txt`. This should cause you to allocate an array of approximately 800 million values. Assuming you're using 4-byte `int`s,

that's 3.2 GB of data. Don't worry about larger cases unless you want to explore alternate approaches. If so, there is a solution which computes the optimal alignment in *linear* space (and quadratic time). This is known as *Hirschberg's algorithm* (1975).

After you have things working, add code to calculate and print execution time. You may use SFML's `sf::Clock` and `sf::Time` classes as follows:

- To your `main.cpp`, add `#include <SFML/System.hpp>` at the top.

- Then define the following object in your `main`

```
sf::Clock clock;
```

- At the end of `main`, after computing the solution, capture the running time:

```
sf::Time t = clock.getElapsedTime();
```

- Then **after** printing out the solution, display the running time:

```
cout << "Execution time is " << t.asSeconds() << " seconds" << endl;
```

3   Valgrind

Verify your algorithm's space usage with the `valgrind` runtime analysis tool.

If necessary, install `valgrind`

```
sudo apt-get install valgrind
```

- Try different compiler optimization flags to see how execution time is affected. For example, `-O1` or `-O2` could halve running time over using no optimization.

- Make sure your code is compiled and linked with debugging information (`-g` flag).

- Use `valgrind` to run your code with the `massif` heap analysis tool. For example:

```
valgrind --tool=massif ./EDistance < sequence/ecoli28284.txt
```

- Then `valgrind` will produce a log file named `massif.out.XXXXX` where `XXXXX` is the process ID from your run. View the file with the `ms_print` utility that is part of the `valgrind` distribution. For example:

```
ms_print massif.out.11515 | less
```

- By naming the `massif.out` file, confirm that the amount of memory used matches your expectations.

- Also, you can use the program `massif-visualizer` to view the logs (see `http://milianw.de/tag/massif-visualizer`. Install with

```
sudo apt-get install massif-visualizer
```

- More documentation on `valgrind` and `massif` is available at `http://valgrind.org/docs/manual/ms-manual.html`.

4   Making the Autograder Happy

The autograder will compile your code and check it with the following input files:

```
bothgaps20.txt    ecoli2500.txt    ecoli7000.txt    example10.txt    fli8.txt
ecoli10000.txt    ecoli5000.txt    endgaps7.txt     fli10.txt        fli9.txt
```

To make sure this works:

- Remember to add `-lfsml-system` to your `LIB` variable in your `Makefile`

- Your `Makefile` must produce an executable named `EDistance`

- Your program must accept input from `stdin` (e.g. using the < redirect).

- Your program must print out **first** the edit distance and **then** the edit path.

- You must follow the format (see example above) exactly, and you can't have trailing spaces at the end of the lines.

- You should print the elapsed time after all this.

- Note that the line endings in the files are not consistent. Some use Unix-style (`\n`) line endings and some use Windows-style (`\r\n`) line endings, and some files do not have final newlines. The line ending is not considered part of the string.

- Everything must be in a directory named ps5, and per usual, make sure to remove object files before tarring up your work.

5   Reporting your results

Please fill out your data in the `Readme-ps5.md` file.

- CPU speed in MHz

- Input size is the length of the problem string for at least these test problems

  - `ecoli2500.txt`

  - `ecoli5000.txt`

  - `ecoli7000.txt`

  - `ecoli10000.txt`

  - `ecoli20000.txt`

  - `ecoli28284.txt`

- Time-to-solution is in seconds

- Memory used is in MB

- Implementation is one of: recursive, recursive-with-memoization, Needleman-Wunsch, Hirschberg, other

- Array method: e.g. vectors, c-arrays, hash-table, other

- Operating system: e.g. x86-unix-native, mac-os-x, ubuntu

- CPU type: e.g. core i3, core i7, AMD, etc.

6   Debugging and Testing

You must write a `test.cpp` that tests all of the public functions you wrote as part of this assignment. You should also use the Boost functions `BOOST_REQUIRE_THROW` and `BOOST_REQUIRE_NO_THROW` to verify that your code properly throws exceptions when appropriate and does not throw an exception when it shouldn't.

Remember that some inputs may have multiple alignments with equal costs. Any of these paths is considered valid. Your tests will be used on two valid instructor implementations that may generate different paths and you must accept both while rejecting each of the broken implementations. Each implementation follows the prescribed specification with no additional public members.

7   Extra Credit

For extra credit, compare and discuss using your own `min3` function with `std::min`, as well as the different optimization levels (`-O0`, `-O1`, `-O2`, `-O3`). Additionally, improving your

implementation so that it can solve larger files (likely using Hirschberg's algorithm) and documenting what you did in the readme is also worth extra credit.

8   What to turn in

Your makefile should build programs named `EDistance` and `test` and a static library `EDistance.a`.

Submit a zip archive to Gradescope containing:

- Your `EDistance` class (`EDistance.cpp` and `EDistance.hpp`)
- Your `main.cpp` file
- Your unit tests (`test.cpp`)
- The makefile for your project. The makefile should have targets `all`, `EDistance`, `test`, `lint` and `clean`. Make sure that all prerequisites are correct.
- Your `Readme-ps5.md` that includes
    1. Your name
    2. Statement of functionality of your program (e.g. fully works, partial functionality, extra credit)
    3. Experimental results as described in Section 5 above.
    4. Key features or algorithms used
    5. Any other notes
- Any other source files that you created.

Make sure that all of your files are in a directory named `ps5` before archiving it and that there are no `.o` or other compiled files in it.

9 Grading rubric

Feature	Points	Comment
Autograder	30	Full & Correct Implementation
	2	Contains all code files and builds
	4	Passes your tests
	2	`min3` and `penalty` static functions
	10	`optDistance`
	10	`alignment`
Analysis	10	
	4	Completes time and space tables
	4	Valgrind analysis
	1	Largest n memory calculation
	1	Largest n time calculation
Readme	5	Complete
	1	Describes the data structure used.
	2	Describes the algorithm used.
	2	Describes the testing decisions.
Extra Credit	8	
	+2	Compares custom `min3` vs `std::min`
	+2	Compares different optimization levels
	+4	Able to solve extra large data sets
Penalties		
	-5	Memory leaks or issues
	-5	Linting problems
	-3	Non-private fields
	-10%	Each day late
Total	45	